

SaaS Subscription Analytics: Revenue Growth, Churn, and Retention Analysis

Project Overview

The goal of this project is to analyze a subscription-based SaaS dataset to understand revenue trends, customer churn behavior, and retention patterns.

Key business questions:

- How is revenue growing over time?
- How many customers are churning?
- Which subscription plans have the highest churn?
- How long do customers stay?
- Are newer customers retaining better than older ones?

Dataset Description

The dataset contains subscription-level SaaS customer data, including:

- `subscription_id`: Unique identifier for each subscription
- `account_id`: Unique customer identifier
- `start_date`: Subscription start date
- `end_date`: Subscription end date (if churned)
- `plan_tier`: Customer subscription plan (Basic, Pro, Enterprise)
- `mrr_amount`: Monthly recurring revenue associated with the subscription
- `churn_flag`: Indicates whether the customer churned

This dataset supports analysis of revenue trends, customer churn, retention patterns, customer lifetime, and cohort behavior.

```
In [2]: import sqlite3
import pandas as pd
```

```
In [3]: import pandas as pd
import sqlite3
import matplotlib.pyplot as plt
```

```
In [4]: df = pd.read_csv('ravenstack_subscriptions.csv')
```

```
In [6]: import sqlite3
import pandas as pd

conn = sqlite3.connect(':memory:')
```

```
df = pd.read_csv('ravenstack_subscriptions.csv')

df.to_sql('subscriptions', conn, index=False, if_exists='replace')
```

Out[6]: 5000

```
In [28]: query = """
select *
from subscriptions
limit 5
"""
pd.read_sql(query, conn)
```

Out[28]:

	subscription_id	account_id	start_date	end_date	plan_tier	seats	mrr_amount	arr_amount
0	S-8cec59	A-3c1a3f	2023-12-23	2024-04-12	Enterprise	14	2786	33432
1	S-0f6f44	A-9b9fe9	2024-06-11	NaN	Pro	17	833	9996
2	S-51c0d1	A-659280	2024-11-25	NaN	Enterprise	62	0	0
3	S-f81687	A-e7a1e2	2024-11-23	2024-12-13	Enterprise	5	995	11940
4	S-cff5a2	A-ba6516	2024-01-10	NaN	Enterprise	27	5373	64476

New Revenue by Signup Month

This analysis calculates monthly recurring revenue based on subscription start dates to understand revenue growth trends over time.

```
In [8]: #New Revenue by Signup Month

query = """
SELECT
    strftime('%Y-%m', start_date) AS month,
    SUM(mrr_amount) AS total_mrr
FROM subscriptions
GROUP BY month
ORDER BY month;
"""
pd.read_sql(query, conn)
```

Out[8]:

	month	total_mrr
0	2023-01	4684
1	2023-02	11079
2	2023-03	25885
3	2023-04	41788

	month	total_mrr
4	2023-05	85919
5	2023-06	74987
6	2023-07	120422
7	2023-08	165621
8	2023-09	116222
9	2023-10	182518
10	2023-11	217219
11	2023-12	267264
12	2024-01	276933
13	2024-02	365442
14	2024-03	421916
15	2024-04	445858
16	2024-05	632549
17	2024-06	537758
18	2024-07	707816
19	2024-08	648426
20	2024-09	992366
21	2024-10	1173608
22	2024-11	1549040
23	2024-12	2273427

This analysis calculates revenue generated from new customer subscriptions based on signup date. It provides insight into customer acquisition trends and revenue growth over time.

In [29]:

```
#New Revenue by Signup Month Visual


query_mrr = """
SELECT
    strftime('%Y-%m', start_date) AS month,
    SUM(mrr_amount) AS total_mrr
FROM subscriptions
GROUP BY month
ORDER BY month;
"""

df_mrr = pd.read_sql(query_mrr, conn)

plt.figure()
plt.plot(df_mrr['month'], df_mrr['total_mrr'])
plt.xticks(rotation=45)

plt.title("New Revenue by Signup Month")
plt.xlabel("Month")
plt.ylabel("New Customer Revenue")
```

```
plt.show()
```



Revenue Growth Insights

- New customer revenue shows strong upward growth over time, increasing significantly throughout 2024
- Growth accelerates in the second half of 2024, indicating increased customer acquisition or customer expansion
- There are some fluctuations month-to-month, but the overall trend reflects rapid business growth

Churn Analysis

This section analyzes how many customers cancel their subscriptions each month.

```
In [12]: #Churn Analysis

query = """
SELECT
    strftime('%Y-%m', end_date) AS churn_month,
    COUNT(*) AS churned_customers
```

```

FROM subscriptions
WHERE churn_flag = 1
      AND end_date IS NOT NULL
GROUP BY churn_month
ORDER BY churn_month;
"""
pd.read_sql(query, conn)

```

Out[12]:

	churn_month	churned_customers
0	2023-04	1
1	2023-06	1
2	2023-07	1
3	2023-08	1
4	2023-10	3
5	2023-11	8
6	2023-12	3
7	2024-01	11
8	2024-02	7
9	2024-03	9
10	2024-04	12
11	2024-05	12
12	2024-06	13
13	2024-07	17
14	2024-08	28
15	2024-09	35
16	2024-10	58
17	2024-11	73
18	2024-12	193

In [13]:


```

#Monthly Churn Trend

df_churn = pd.read_sql(query, conn)

plt.figure()
plt.plot(df_churn['churn_month'], df_churn['churned_customers'])
plt.xticks(rotation=45)
plt.title("Monthly Churned Customers")
plt.xlabel("Month")
plt.ylabel("Churned Customers")
plt.show()

```



Monthly Churn Trend Insights

- Churn increases alongside growth, suggesting onboarding and customer success processes are not scaling effectively. The company should investigate onboarding experience and support capacity during high-growth periods
- The number of churned customers rises alongside overall growth, suggesting scaling challenges
- The large spike in December 2024 may indicate retention issues or seasonal behavior

In [14]:

```
#Active Customers Per Month

query_active = """
SELECT
    strftime('%Y-%m', start_date) AS month,
    COUNT(DISTINCT account_id) AS active_customers
FROM subscriptions
GROUP BY month
ORDER BY month;
"""

df_active = pd.read_sql(query_active, conn)
df_active
```

Out [14]:

	month	active_customers
0	2023-01	2
1	2023-02	9
2	2023-03	16
3	2023-04	25
4	2023-05	26
5	2023-06	36
6	2023-07	45
7	2023-08	63
8	2023-09	51
9	2023-10	65
10	2023-11	86
11	2023-12	93
12	2024-01	103
13	2024-02	111
14	2024-03	146
15	2024-04	137
16	2024-05	168
17	2024-06	168
18	2024-07	217
19	2024-08	212
20	2024-09	248
21	2024-10	273
22	2024-11	300
23	2024-12	335

In [15]:

#Reuse Churn Data

```

query_churn = """
SELECT
    strftime('%Y-%m', start_date) AS month,
    COUNT(*) AS churned_customers
FROM subscriptions
WHERE churn_flag = 1
    AND end_date IS NOT NULL
GROUP BY month
ORDER BY month;
"""

```

```

df_churn = pd.read_sql(query_churn, conn)
df_churn

```

Out[15]:

	month	churned_customers
0	2023-02	4
1	2023-03	2
2	2023-04	2
3	2023-05	5
4	2023-06	4
5	2023-07	10
6	2023-08	8
7	2023-09	14
8	2023-10	9
9	2023-11	16
10	2023-12	11
11	2024-01	14
12	2024-02	19
13	2024-03	11
14	2024-04	19
15	2024-05	21
16	2024-06	19
17	2024-07	21
18	2024-08	26
19	2024-09	43
20	2024-10	42
21	2024-11	63
22	2024-12	103

Churn Rate

Churn rate is calculated as the percentage of customers who cancel in a given month relative to total customers.

In [16]:

```
#Churn Rate Calculation
```

```
df_merged = pd.merge(df_active, df_churn, on='month', how='left')
df_merged['churned_customers'] = df_merged['churned_customers'].fillna(0)

df_merged["churn_rate"] = df_merged["churned_customers"] / df_merged["active_cus
df_merged
```


Out[16]:

	month	active_customers	churned_customers	churn_rate
0	2023-01	2	0.0	0.000000
1	2023-02	9	4.0	0.444444
2	2023-03	16	2.0	0.125000
3	2023-04	25	2.0	0.080000
4	2023-05	26	5.0	0.192308
5	2023-06	36	4.0	0.111111
6	2023-07	45	10.0	0.222222
7	2023-08	63	8.0	0.126984
8	2023-09	51	14.0	0.274510
9	2023-10	65	9.0	0.138462
10	2023-11	86	16.0	0.186047
11	2023-12	93	11.0	0.118280
12	2024-01	103	14.0	0.135922
13	2024-02	111	19.0	0.171171
14	2024-03	146	11.0	0.075342
15	2024-04	137	19.0	0.138686
16	2024-05	168	21.0	0.125000
17	2024-06	168	19.0	0.113095
18	2024-07	217	21.0	0.096774
19	2024-08	212	26.0	0.122642
20	2024-09	248	43.0	0.173387
21	2024-10	273	42.0	0.153846
22	2024-11	300	63.0	0.210000
23	2024-12	335	103.0	0.307463

In [17]:

#Churn Rate Calculation Visual

```
plt.figure()
plt.plot(df_merged['month'], df_merged['churn_rate'])
plt.xticks(rotation=45)
plt.title("Monthly Churn Rate")
plt.xlabel("Month")
plt.ylabel("Churn Rate")
plt.show()
```



Monthly Churn Rate Insights

- Churn rate fluctuates over time, with higher volatility in early months due to a smaller customer base
- As the customer base grows, churn rate becomes more stable but increases again in later months
- The spike in churn rate toward the end of the period suggests potential retention or product issues

In [18]:

#Churn Count by Plan

```
query = """
SELECT
    plan_tier,
    COUNT(*) AS churned_customers
FROM subscriptions
WHERE churn_flag = 1
GROUP BY plan_tier
ORDER BY churned_customers DESC;
"""
pd.read_sql(query, conn)
```


Out[18]:

	plan_tier	churned_customers
0	Enterprise	172
1	Pro	162
2	Basic	152

```
In [19]: #Churn Count by Plan Plan Visual

df_plan_churn = pd.read_sql(query, conn)

plt.figure()
plt.bar(df_plan_churn['plan_tier'], df_plan_churn['churned_customers'])
plt.title("Churn by Plan Tier")
plt.xlabel("Plan")
plt.ylabel("Churned Customers")
plt.show()
```



Churn by Tier Insights

- The Enterprise plan has the highest total churn, followed closely by Pro and Basic plans
- Higher churn in Enterprise may be due to a larger customer base or higher expectations from customers
- Differences across plans suggest varying customer behavior depending on pricing or features

Churn by Plan

This analysis identifies which subscription plans have the highest churn.

In [20]:

```
#Churn Rate by Plan

query = """
SELECT
    plan_tier,
    COUNT(*) AS total_customers,
    SUM(CASE WHEN churn_flag = 1 THEN 1 ELSE 0 END) AS churned_customers,
    ROUND(
        1.0 * SUM(CASE WHEN churn_flag = 1 THEN 1 ELSE 0 END) / COUNT(*), 3) AS churn_rate
FROM subscriptions
GROUP BY plan_tier
ORDER BY churn_rate DESC;
"""
pd.read_sql(query, conn)
```

Out[20]:


	plan_tier	total_customers	churned_customers	churn_rate
0	Enterprise	1723	172	0.100
1	Pro	1675	162	0.097
2	Basic	1602	152	0.095

In [21]:

```
#Churn Rate by Plan Visual

df_plan_rate = pd.read_sql(query, conn)

plt.figure()
plt.bar(df_plan_rate['plan_tier'], df_plan_rate['churn_rate'])
plt.title("Churn Rate by Plan")
plt.xlabel("Plan")
plt.ylabel("Churn Rate")
plt.show()
```



Churn Rate Insights

- The Enterprise plan has the highest churn rate, indicating relatively weaker retention compared to other plans
- Churn rates across all plans are fairly close, suggesting consistent retention challenges across tiers
- Slightly higher churn in Enterprise may indicate pricing or value expectations are not fully met

Retention Analysis

This section analyzes how long customers stay subscribed before churning.

```
In [22]: #Customer Lifetime by Plan (Monthly)

query = """
SELECT
    subscription_id,
    plan_tier,
    start_date,
    end_date,
    CASE
        WHEN end_date IS NOT NULL
        THEN (julianday(end_date) - julianday(start_date)) / 30.0
        ELSE NULL
    
```

```
END AS lifetime_months
FROM subscriptions
WHERE churn_flag = 1;
"""

df_lifetime = pd.read_sql(query, conn)
df_lifetime.head()
```

Out [22]:

	subscription_id	plan_tier	start_date	end_date	lifetime_months
0	S-8cec59	Enterprise	2023-12-23	2024-04-12	3.700000
1	S-f81687	Enterprise	2024-11-23	2024-12-13	0.666667
2	S-13939b	Pro	2023-07-21	2024-08-03	12.633333
3	S-9333c5	Basic	2024-06-06	2024-08-13	2.266667
4	S-3bd427	Basic	2024-12-04	2024-12-27	0.766667

In [23]:

```
#Average Customer Lifetime by Plan (Monthly)

query = """
SELECT
    plan_tier,
    ROUND(AVG((julianday(end_date) - julianday(start_date)) / 30.0), 2) AS avg_lif
FROM subscriptions
WHERE churn_flag = 1
GROUP BY plan_tier
ORDER BY avg_lifetime_months DESC;
"""

pd.read_sql(query, conn)
```

Out [23]:


	plan_tier	avg_lifetime_months
0	Basic	3.34
1	Enterprise	2.79
2	Pro	2.72

In [24]:

```
#Average Customer Lifetime by Plan (Monthly) Visual

df_lifetime_tier = pd.read_sql(query, conn)

plt.figure()
plt.bar(df_lifetime_tier['plan_tier'], df_lifetime_tier['avg_lifetime_months'])
plt.title("Average Customer Lifetime by Plan")
plt.xlabel("Plan")
plt.ylabel("Months")
plt.show()
```



In []:

Retention Insights

- Customers on the Basic plan have the longest average lifetime, indicating stronger retention at lower price tiers
- Enterprise and Pro plans have slightly shorter lifetimes, suggesting potential challenges in retaining higher-value customers
- Overall customer lifetime is relatively short across all plans, highlighting an opportunity to improve long-term retention

Cohort Analysis

Cohort analysis groups customers by their signup month to evaluate retention trends over time. This analysis groups customers by their signup month (cohort) and calculates the average lifetime of each cohort.

The goal is to understand whether customers who joined at different times behave differently in terms of retention.

In [25]:

```
#Cohort Analysis (Average Customer Lifetime)
```

```

query = """
SELECT
    account_id,
    plan_tier,
    strftime('%Y-%m', start_date) AS cohort_month,
    CASE
        WHEN end_date IS NOT NULL
        THEN (julianday(end_date) - julianday(start_date)) / 30.0
        ELSE (julianday('now') - julianday(start_date)) / 30.0
    END AS lifetime_months
FROM subscriptions;
"""

df_cohort = pd.read_sql(query, conn)
df_cohort.head()

```

Out[25]:

	account_id	plan_tier	cohort_month	lifetime_months
0	A-3c1a3f	Enterprise	2023-12	3.700000
1	A-9b9fe9	Pro	2024-06	23.135621
2	A-659280	Enterprise	2024-11	17.568955
3	A-e7a1e2	Enterprise	2024-11	0.666667
4	A-ba6516	Enterprise	2024-01	28.235621

This table summarizes each cohort (grouped by signup month) by showing the total number of customers and their average lifetime.

It provides a high-level view of how customer retention varies across different signup periods.

In [26]:

```

#Cohort Summary Table

query = """
SELECT
    strftime('%Y-%m', start_date) AS cohort_month,
    COUNT(*) AS total_customers,
    ROUND(AVG(
        CASE
            WHEN end_date IS NOT NULL
            THEN (julianday(end_date) - julianday(start_date)) / 30.0
            ELSE (julianday('now') - julianday(start_date)) / 30.0
        END
    ), 2) AS avg_lifetime
FROM subscriptions
GROUP BY cohort_month
ORDER BY cohort_month;
"""

df_cohort_summary = pd.read_sql(query, conn)

print(df_cohort_summary.to_string(index=False))

```


cohort_month	total_customers	avg_lifetime
2023-01	3	40.19
2023-02	11	28.22
2023-03	17	35.74
2023-04	31	35.18

2023-05	29	31.94
2023-06	46	33.20
2023-07	59	29.81
2023-08	82	30.71
2023-09	62	26.40
2023-10	84	28.40
2023-11	121	27.12
2023-12	121	27.00
2024-01	134	25.57
2024-02	152	24.22
2024-03	186	24.81
2024-04	193	22.97
2024-05	245	22.34
2024-06	248	21.48
2024-07	330	20.75
2024-08	337	19.46
2024-09	440	18.11
2024-10	504	17.41
2024-11	612	16.08
2024-12	953	15.04

In [27]:

```
#Cohort Summary Visual
```

```
plt.figure()
plt.plot(df_cohort_summary['cohort_month'], df_cohort_summary['avg_lifetime'])
plt.xticks(rotation=45)
plt.title("Average Customer Lifetime by Cohort")
plt.xlabel("Cohort Month")
plt.ylabel("Avg Lifetime (Months)")
plt.show()
```



Cohort Insights

- Earlier cohorts show significantly higher average lifetime compared to more recent cohorts
- Average lifetime declines over time, suggesting newer customers are not retaining as long
- This may indicate changes in customer quality, onboarding experience, or product-market fit over time

Key Findings

- New customer revenue shows strong and accelerating growth, particularly in late 2024
- Churn increases alongside growth, indicating potential retention challenges at scale
- Churn rate stabilizes but rises again in later months, suggesting emerging issues
- Enterprise plan has the highest churn rate, pointing to possible pricing or expectation gaps
- Customer lifetime is relatively short across all plans, highlighting retention opportunities
- Cohort analysis suggests that newer customers are retaining less effectively than earlier cohorts

Actionable Recommendations

- Improve onboarding for new customers (cohort decline)
- Investigate Enterprise churn drivers (pricing / expectations gap)
- Introduce retention campaigns before month 3 (short lifetime)
- Monitor churn spikes in Q4 (seasonality risk)